

LAABH GUPTA

Full-Stack Developer | AI/ML Engineer | Data Analyst

PERSONAL INFORMATION

Phone : +91 9793084444

Email : reachlaabhgupta@gmail.com

LinkedIn : [linkedin.com/in/laabhgupta](https://www.linkedin.com/in/laabhgupta)

GitHub : github.com/Laabh-Gupta

LeetCode : leetcode.com/u/laabhgupta

HackerRank : hackerrank.com/profile/laabhgupta

SUMMARY

Results-oriented developer with a B.Tech in Computer Science (AI & ML specialization) and hands-on industry experience in agentic AI and data engineering at EY GDS. Skilled in designing, building, and deploying scalable full-stack applications, machine learning models, and enterprise data pipelines. Strong foundation in Python, Java, and data structures combined with expertise in React.js, Node.js, Databricks, and cloud platforms. Proven ability to deliver measurable impact—improving decision-making speed, streamlining workflows, and enhancing user experience. Eager to apply data-driven problem-solving to innovative projects in software engineering, AI/ML, and analytics.

CORE SKILLS

Languages: Java, Python, C, C++, SQL

AI/ML: MLflow, Agentic AI, RAG, PyTorch, TensorFlow, Scikit-learn, OpenCV, Pandas, NumPy

Frontend: Streamlit, HTML, CSS, JavaScript, TypeScript, React.js

Backend: Spring Boot, Node.js, FastAPI, REST APIs, MVC Architecture, Spring Security (JWT), Uvicorn

Data & Cloud: Databricks, MySQL, MongoDB, PostgreSQL

Tools & Platforms: Git, GitHub, Docker, CI/CD, Jupyter, VS Code, IntelliJ

Methodologies: Agile, Scrum, Waterfall

EDUCATION

B.Tech in Computer Science Engineering with Artificial Intelligence and Machine Learning
Vellore Institute of Technology, Chennai | 2022-26 | CGPA 8.1

Secondary School

St. Thomas School, Kanpur | 2021-22 | 92%

High School

St. Thomas School, Kanpur | 2019-20 | 89%

EXPERIENCE

Ernst & Young Global Delivery Services

AI & Data Intern (Jan 2026 – May 2026)

Worked as an AI & Data Engineer Intern, contributing to two enterprise-scale projects in the data engineering and MLOps space, while upskilling in Generative AI, RAG, Agentic AI, Databricks, and MLOps through structured learning and hands-on implementation.

Project 1 — Agentic Data Ingestion Platform One liner: An agentic AI system that automates discovery and ingestion of enterprise data from multiple cloud sources into a unified platform.

- **Business Need:** Enterprises managing data across disparate sources faced slow, manual, and error-prone ingestion pipelines with no standardization or scalability.
- **Business Benefits:** Faster ingestion setup, reduced errors, lower dependency on experienced data engineers, reusable framework scalable across multiple clients.
- **Tech Stack:** Streamlit, Databricks, Databricks Unity Catalog, LangChain, REST APIs, JSON Config Framework, Control Plane Service.

Project 2 — MLOps Experiment Tracking Dashboard One liner: A full-stack MLOps dashboard for real-time visibility into ML experiments, model runs, and performance metrics across a Databricks workspace.

- **Business Need:** Teams lacked a centralized view to track MLflow experiments and compare model runs; default MLflow UI had limited customization for stakeholders.
- **Business Benefits:** Centralized experiment visibility, faster model comparison, reduced manual tracking, better governance and traceability, lower risk of faulty deployments.
- **Tech Stack:** Databricks MLflow, Databricks SQL / Delta Tables, PySpark, FastAPI, Uvicorn, React, Axios, Python ML Libraries (scikit-learn, pandas, numpy).

Hindustan Aeronautics Limited

Tech Intern(June 2025 – July 2025)

Developed and deployed a secure full-stack application to automate cost estimation for aircraft parts manufacturing.

- Built a modular Node.js + MySQL backend system for retrieving, computing, and reporting man-hour, man standard rate, and material costs from internal production databases.
- Integrated Express-session-based user authentication using bcrypt for password hashing, enabling secure and role-based access control for internal engineering teams.
- Improved decision-making speed by ~25% by streamlining computation pipelines using MVC architecture.
- Leveraged Agile methodology and CI/CD pipelines with GitHub for automated version control and deployment.

PROJECTS

Voice Anti-Spoofing System

Project Lead (Academic Project) | December 2024 – March 2025

Deep learning system to differentiate AI-generated voices from human speech, deployed as a full-stack web application.

- Engineered an end-to-end pipeline, transforming audio to Mel Spectrograms and fine-tuning multiple architectures (CNNs, Vision Transformer) with data augmentation to achieve 99.75% test accuracy in spoof detection.
- Developed and deployed two full-stack web applications using a FastAPI backend to serve PyTorch models via a REST API, and a React.js frontend for real-time user interaction and audio classification.
- Prior research in voice anti-spoofing primarily utilizes standard CNNs on spectrograms, often without demonstrating a complete, deployed application for end-user interaction.
- Innovated by successfully adapting a Vision Transformer (ViT) for superior audio classification and bridged the research-to-production gap by deploying the final models into fully functional web applications.
- To extend the system to support real-time audio stream analysis from microphones and investigate specialized audio transformer architectures for enhanced robustness against diverse spoofing attacks.

Daily Journal App

A full-stack journaling platform with a secure Spring Boot backend and modern UI, enabling users to create, edit, and manage daily entries with persistent data storage.

- Engineered a RESTful API using Spring Boot, Spring Data JPA, and MongoDB, enabling robust CRUD operations for journal entries.
- Implemented Spring Security with JWT authentication to ensure secure, role-based access control.
- Optimized database interactions with Hibernate ORM, ensuring efficient persistence and retrieval.
- Deployed locally with Maven + Node.js, demonstrating end-to-end full-stack development and integration.
- Integrated Confluent Kafka to enable asynchronous data transfer and event-driven communication, improving scalability and system responsiveness.

ACHIEVEMENTS

- Solved 370+ coding problems on LeetCode, ranking in the top 2% globally, demonstrating advanced problem-solving and algorithmic skills.
- Completed the Databricks Machine Learning Practitioner Learning Plan — all constituent courses including PySpark, MLflow, ML model lifecycle, and MLOps (certified during EY GDS internship).
- Completed '[The Complete Full-Stack Web Development Bootcamp](#)' a comprehensive 61.5-hour course on Udemy, mastering key front-end and back-end technologies.
- Achieved certification in '[Data Structure & Algorithms in Java](#)' from Great Learning Academy, reinforcing advanced problem-solving abilities.